

Convolutional Neural Network






Image Classification

- A core task in Computer Vision



Challenges of Recognition

Viewpoint




Illumination



This image is [CC0 1.0 public domain](#)

Deformation



This image by [Umberto Salvagnin](#) is licensed under [CC-BY 2.0](#)

Occlusion



This image by [jonsson](#) is licensed under [CC-BY 2.0](#)

Clutter



This image is [CC0 1.0 public domain](#)


Intraclass Variation



This image is [CC0 1.0 public domain](#)

Convolutional Neural Network

- Most widely used for image classification.
- Generally, it consists of convolution layer, pooling layer and fully-connected layer.
- Convolution, Pooling layer – feature extraction
- Fully-connected layer – classification




Convolution Filters(Hand Crafted)



0	0	0	0	0
0	0	-1	0	0
0	-1	5	-1	0
0	0	-1	0	0
0	0	0	0	0



	0	1	0	
	1	-4	1	
	0	1	0	



0	0	0	0	0
0	1	1	1	0
0	1	1	1	0
0	1	1	1	0
0	0	0	0	0



	-2	-1	0	
	-1	1	1	
	0	1	2	



Let's Try!

- <http://setosa.io/ev/image-kernels/>


0	-1	0
-1	5	-1
0	-1	0

sharpen ▼




CNN 동작원리

- 이미지를 작은 tile로 나누고, convolution filter를 통해 tile에서 특정 feature를 추출(예: 귀)
- Filter가 다음 tile로 이동하면서 같은 방법으로 feature를 추출(동일한 weight 사용)
- 다른 feature(예: 눈)를 추출하는 filter를 추가로 만들고 위와 같은 방법으로 tile을 하나씩 network에 적용
- 추출된 모든 feature들을 잘 조합하여 최종적으로 이미지를 판단



2D Convolution Layer


Convolution Layer




2D Convolution Layer

Convolution Layer

consider a second, **green** filter




Feature maps
Activation maps



2D Convolution Layer


For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:



We stack these up to get a “new image” of size 28x28x6!


Dense Layer vs 1-D Convolution Layer

- Dense Layer(Fully Connected Layer)
 - $y_0 = x_0 \cdot w_{00} + x_1 \cdot w_{01} + x_2 \cdot w_{02} + x_3 \cdot w_{03}$



Dense Layer vs 1-D Convolution Layer


- Dense Layer(Fully Connected Layer)
 - $y_0 = x_0 \cdot w_{00} + x_1 \cdot w_{01} + x_2 \cdot w_{02} + x_3 \cdot w_{03}$
 - $y_1 = x_0 \cdot w_{10} + x_1 \cdot w_{11} + x_2 \cdot w_{12} + x_3 \cdot w_{13}$



Dense Layer vs 1-D Convolution Layer


- 1-D Convolution Layer

- $y_0 = x_0 \cdot w_0 + x_1 \cdot w_1 + x_2 \cdot w_2$



Dense Layer vs 1-D Convolution Layer


- 1-D Convolution Layer
 - $y_0 = x_0 \cdot w_0 + x_1 \cdot w_1 + x_2 \cdot w_2$
 - $y_0 = x_1 \cdot w_0 + x_2 \cdot w_1 + x_3 \cdot w_2$



Dense Layer vs 1-D Convolution Layer


- 1-D Convolution Layer

- $y_0 = x_0 \cdot w_0 + x_1 \cdot w_1 + x_2 \cdot w_2$
 - $y_0 = x_1 \cdot w_0 + x_2 \cdot w_1 + x_3 \cdot w_2$
- Weight sharing
&
Locally connected




2D Convolution Layer – Computation

- $1 \times 1 + 1 \times 0 + 1 \times 1 + 0 \times 0 + 1 \times 1 + 1 \times 0 + 0 \times 1 + 0 \times 0 + 1 \times 1 = 4$




2D Convolution Layer – Computation

- $1 \times 1 + 1 \times 0 + 0 \times 1 + 1 \times 0 + 1 \times 1 + 1 \times 0 + 0 \times 1 + 1 \times 0 + 1 \times 1 = 3$



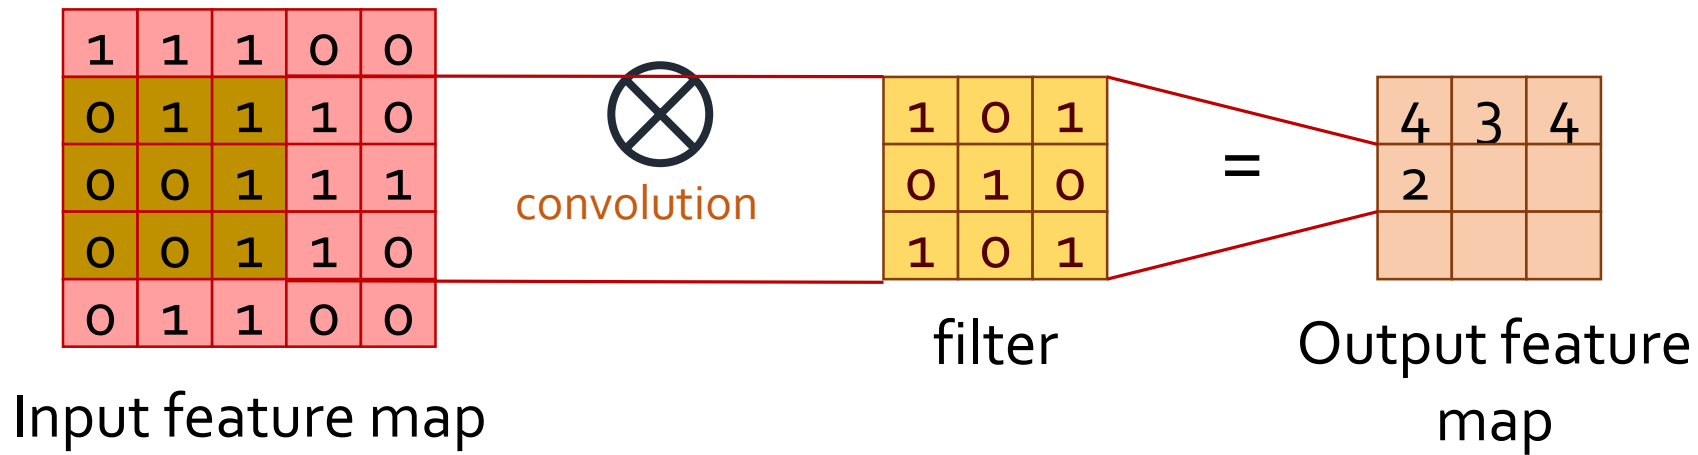
2D Convolution Layer – Computation

- $1 \times 1 + 0 \times 0 + 0 \times 1 + 1 \times 0 + 1 \times 1 + 0 \times 0 + 1 \times 1 + 1 \times 0 + 1 \times 1 = 4$



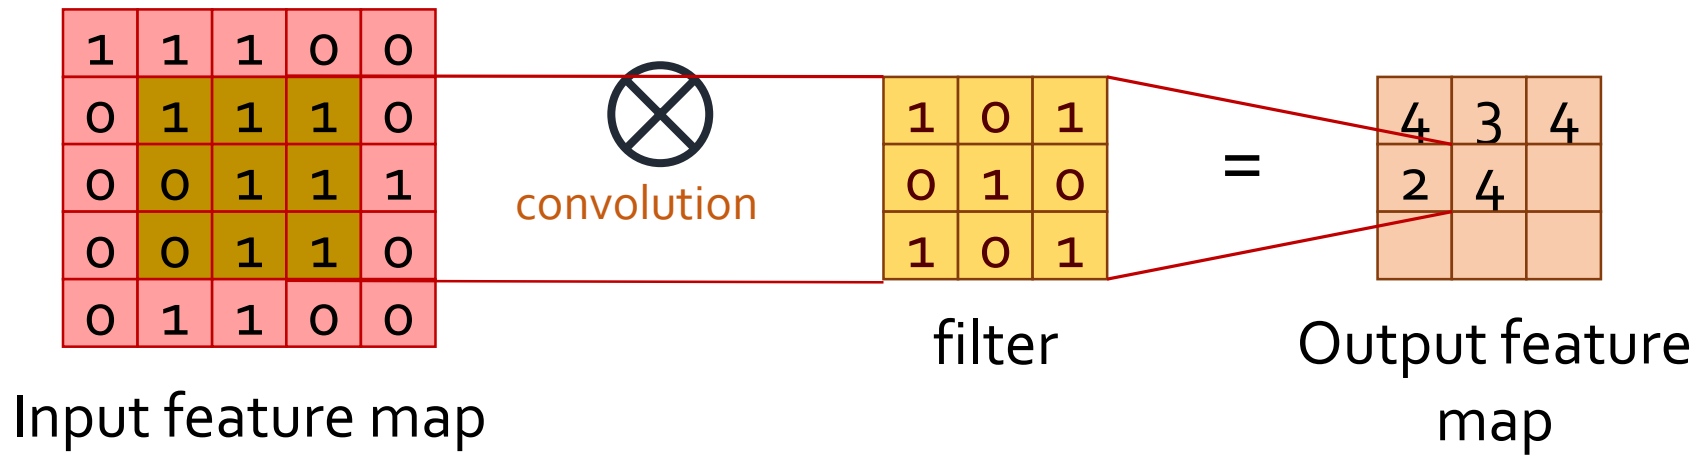
2D Convolution Layer – Computation

- $0 \times 1 + 1 \times 0 + 1 \times 1 + 0 \times 0 + 0 \times 1 + 1 \times 0 + 0 \times 1 + 0 \times 0 + 1 \times 1 = 2$



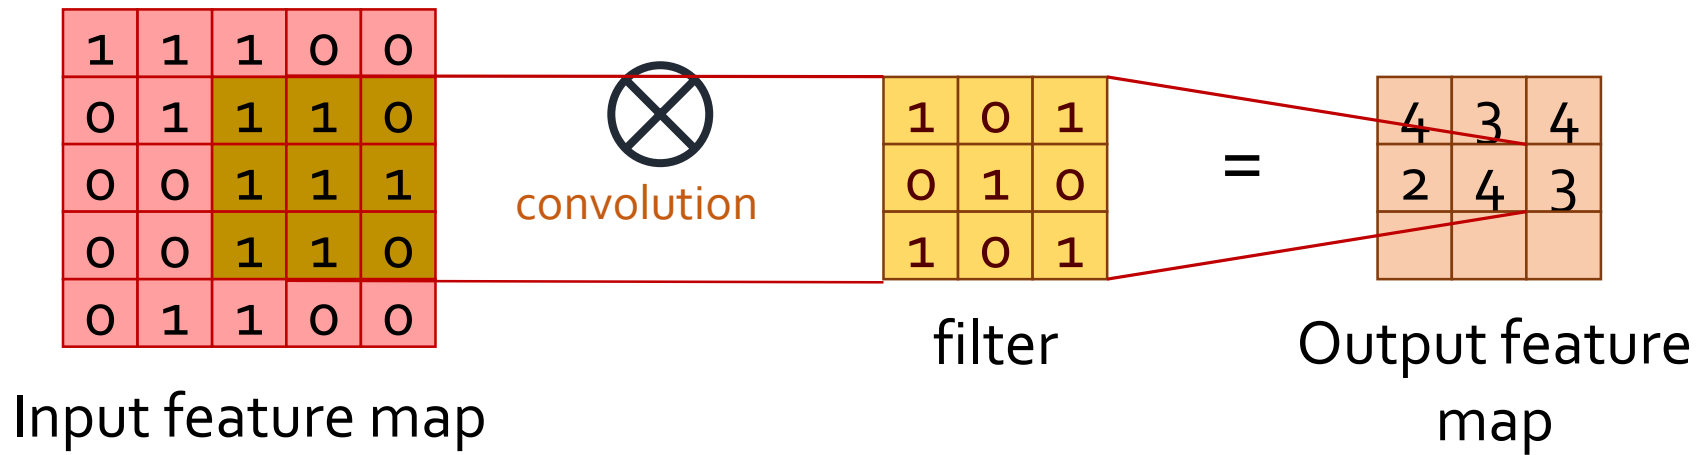
2D Convolution Layer – Computation

- $1 \times 1 + 1 \times 0 + 1 \times 1 + 0 \times 0 + 1 \times 1 + 1 \times 0 + 0 \times 1 + 1 \times 0 + 1 \times 1 = 4$




2D Convolution Layer – Computation

- $1 \times 1 + 1 \times 0 + 0 \times 1 + 1 \times 0 + 1 \times 1 + 1 \times 0 + 1 \times 1 + 1 \times 0 + 0 \times 1 = 3$




2D Convolution Layer – Computation

- $0 \times 1 + 0 \times 0 + 1 \times 1 + 0 \times 0 + 0 \times 1 + 1 \times 0 + 0 \times 1 + 1 \times 0 + 1 \times 1 = 2$




2D Convolution Layer – Computation

- $0 \times 1 + 1 \times 0 + 1 \times 1 + 0 \times 0 + 1 \times 1 + 1 \times 0 + 1 \times 1 + 1 \times 0 + 0 \times 1 = 3$



2D Convolution Layer – Computation

- $1 \times 1 + 1 \times 0 + 1 \times 1 + 1 \times 0 + 1 \times 1 + 0 \times 0 + 1 \times 1 + 0 \times 0 + 0 \times 1 = 4$



2D Convolution Layer – Computation

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image


4		

Convolved
Feature

Feature Extractor

0	0	0	0	0	30	0
0	0	0	0	30	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	0	0	0	0


Pixel representation of filter



Visualization of a curve detector filter



Original image



Visualization of the filter on the image

Feature Extractor



Visualization of the receptive field

0	0	0	0	0	0	30
0	0	0	0	50	50	50
0	0	0	20	50	0	0
0	0	0	50	50	0	0
0	0	0	50	50	0	0
0	0	0	50	50	0	0
0	0	0	50	50	0	0

Pixel representation of the receptive field

*

0	0	0	0	0	30	0
0	0	0	0	30	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	0	0	0	0

Pixel representation of filter

Multiplication and Summation = $(50*30)+(50*30)+(50*30)+(20*30)+(50*30) = 6600$ (A large number!)



Visualization of the filter on the image

0	0	0	0	0	0	0
0	40	0	0	0	0	0
40	0	40	0	0	0	0
40	20	0	0	0	0	0
0	50	0	0	0	0	0
0	0	50	0	0	0	0
25	25	0	50	0	0	0

Pixel representation of receptive field


*

0	0	0	0	0	30	0
0	0	0	0	30	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	0	0	0	0

Pixel representation of filter

Multiplication and Summation = 0

Convolution (Multi Channel, Many Filters)



			1	1	1	0	1	0
		0	1	1	1	0	1	0
1	1	1	0	0				
0	1	1	1	0				
0	0	1	1	1				
0	0	1	1	0				
0	1	1	0	0				

Input channel : 3

\otimes
convolution

				1	0	0
		0	1	0	0	
1	0	1				
0	1	0				
1	0	1				

				1	0	0	1
		0	1	0	0		
1	0	1					
0	-1	0					
1	0	1					

of filters : 2


=

			1	1	1
		3	0	1	
-2	0	2			
0	2	3			


Output channel : 2

Convolution


(Multi Channel, Many Filters)




Visualization of a Convolution Layer



2D Convolution Layer




2D Convolution Layer – 4D Tensors



Options of Convolution

- Stride : filter가 한 번 convolution을 수행 한 후 옆으로(혹은 아래로) 얼마나 이동할 것인가
 - 예) 7x7 input, 3x3 convolution filter with stride 2 \rightarrow 3x3 output!



Options of Convolution

- Zero Padding

0	0	0	0	0	0			
0								
0								
0								
0								

e.g. input 7x7

3x3 filter, applied with **stride 1**

pad with 1 pixel border => what is the output?

7x7 output!

in general, common to see CONV layers with stride 1, filters of size $F \times F$, and zero-padding with $(F-1)/2$. (will preserve size spatially)

e.g. $F = 3 \Rightarrow$ zero pad with 1

$F = 5 \Rightarrow$ zero pad with 2

$F = 7 \Rightarrow$ zero pad with 3

Quiz

- 다음의 각 경우에 convolution layer의 output size는?

1. 32x32x3 input, 10 5x5 filters with stride 1, pad 0
2. 32x32x3 input, 10 5x5 filters with stride 1, pad 2
3. 32x32x3 input, 10 3x3 filters with stride 2, pad 1

- Answer

1. 28x28x10
2. 32x32x10
3. 16x16x10



Input $O/W_i \times H_i \times C_i$ 이고,
 $F \times F$ filter를 K 개 사용하고,
stride는 S ,
zero padding은 P 만큼 했을 경우,
output feature map size($W_o \times H_o \times C_o$)는,

$$W_o = \frac{(W_i - F + 2P)}{S} + 1$$

$$H_o = \frac{(H_i - F + 2P)}{S} + 1$$


$$C_o = K$$

ReLU





Pooling Layer

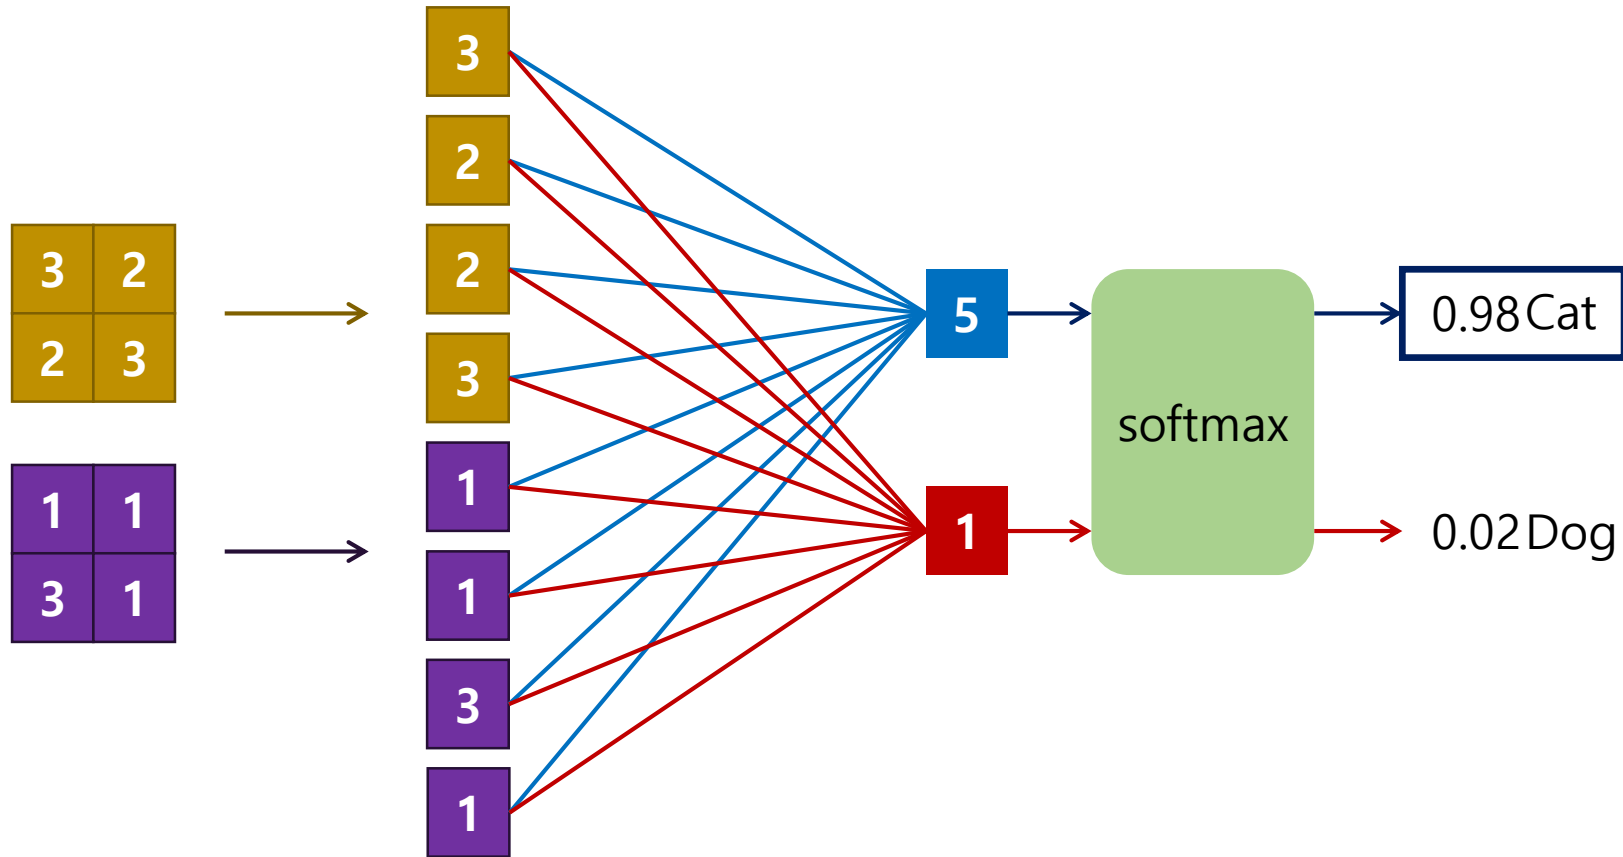
- Max pooling or Average Pooling



2x2 Max Pooling with Stride=1

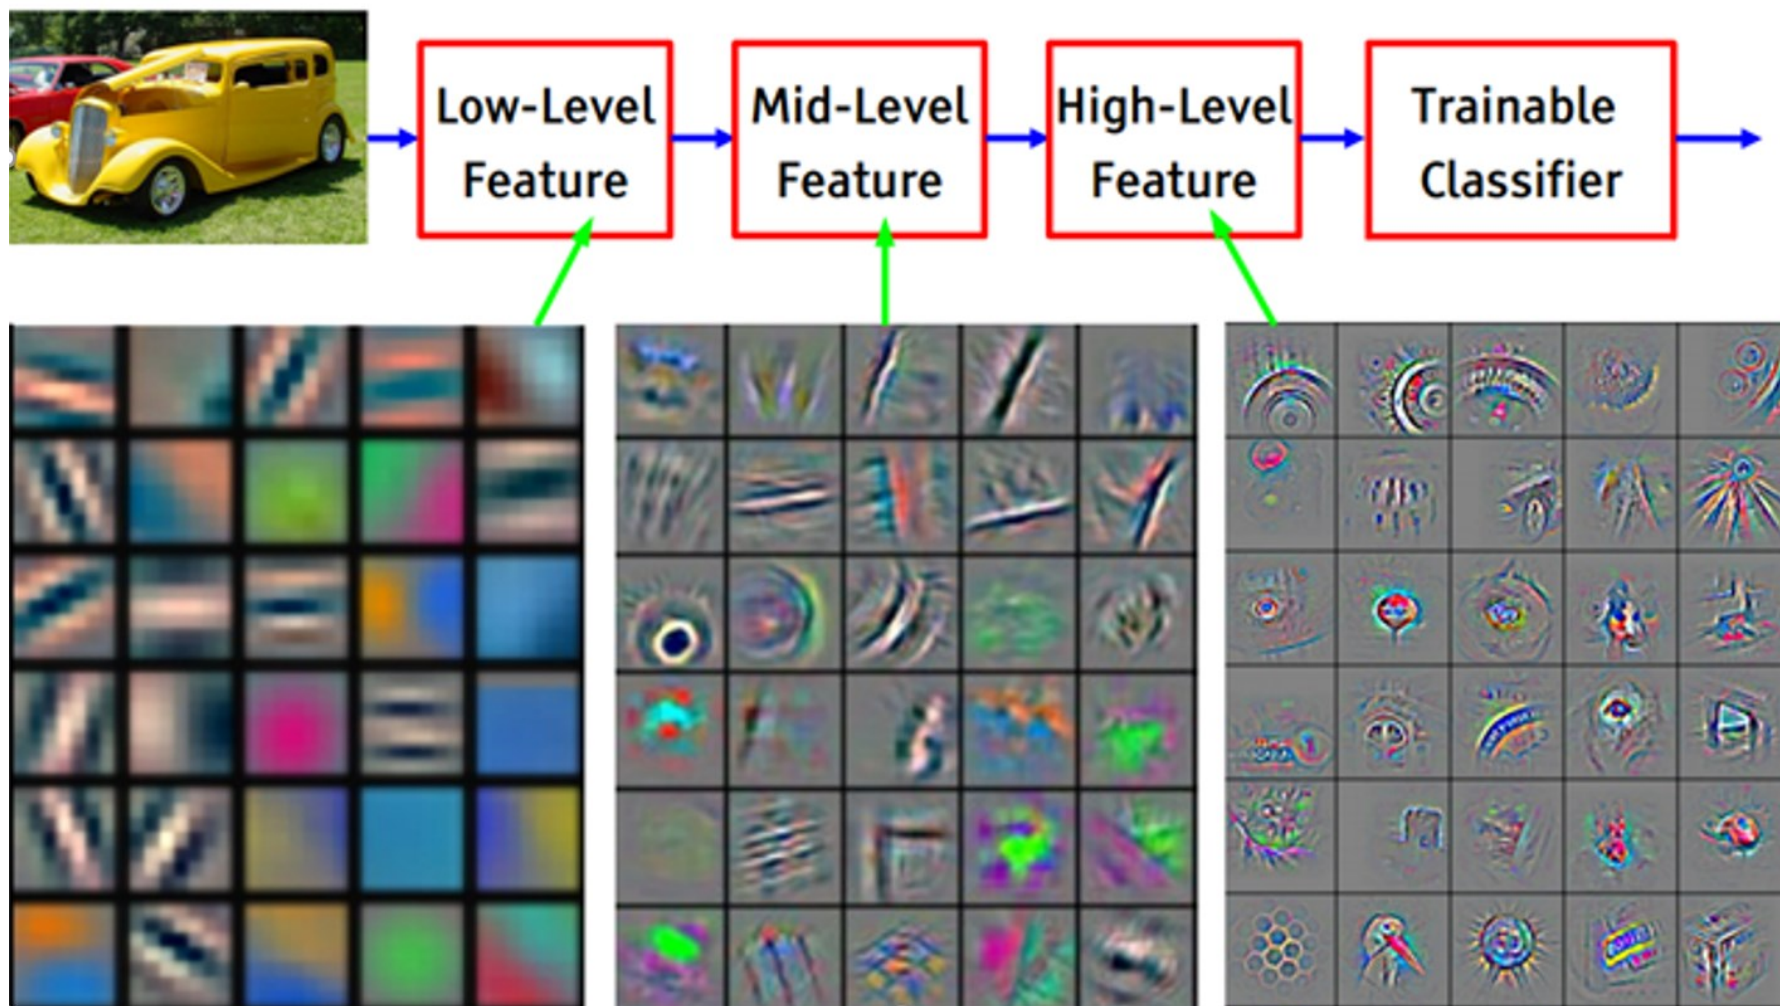


Fully-Connected Layer



Convolutional Neural Network

State of the art object recognition using CNNs



Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

고전적인 CNN의 특징

- Convolution Layer – parameter(weight) sharing
- Good for local invariance – pooling
- 연산량은 Convolution layer가 대부분을 차지
- Parameter 수는 FC layer가 대부분을 차지

Model	Params (M)	Conv (%)	FC (%)	Ops (M)	Conv (%)	FC (%)
AlexNet	61	3.8	96.2	725	91.9	8.1
VGG-F	99	2.2	97.8	762	87.4	12.6
VGG-M	103	6.3	93.7	1678	94.3	5.7
VGG-S	103	6.3	93.7	2640	96.3	3.7
VGG-16	138	10.6	89.4	15484	99.2	0.8
VGG-19	144	13.9	86.1	19647	99.4	0.6
NIN	7.6	100	0	1168	100.0	0.0
GoogLeNet	6.9	85.1	14.9	1566	99.9	0.1